

Convolutional Neural Networks (CNNs)

Ankit Kumar

February 4, 2026

1 Introduction to CNNs

1.1 What is a CNN?

A **Convolutional Neural Network (CNN)** is a specialized deep learning architecture designed for processing **grid-like data**, particularly images. Unlike traditional neural networks that treat input as flat vectors, CNNs preserve spatial relationships.

1.2 Why Not Use Regular Neural Networks (ANNs) for Images?

- **Huge Number of Parameters:** For a 32x32 RGB image:

$$32 \times 32 \times 3 = 3,072 \text{ input neurons}$$

If connected to just 100 neurons in next layer:

$$3,072 \times 100 = 307,200 \text{ weights!}$$

This leads to **computational explosion**.

- **Loss of Spatial Information:** Flattening destroys local patterns (edges, textures).
- **Overfitting:** Too many parameters → poor generalization.

2 CNN Architecture Overview

A typical CNN has this structure:

Input Image → Convolution Layers → Pooling Layers → More Conv/Pool → Flatten → Fully Connected → Output

3 Convolution Layer: The Feature Detector

3.1 What is Convolution?

Imagine sliding a small magnifying glass (filter/kernel) over an image to detect patterns.

3.2 Mathematical Definition

For 2D convolution:

$$(I * K)[i, j] = \sum_m \sum_n I[i - m, j - n] \cdot K[m, n]$$

Where:

- I = Input image (matrix)
- K = Filter/kernel (small matrix)
- $*$ = Convolution operation

3.3 Example Calculation

Let's convolve a 4×4 image with a 3×3 filter:

Image I :

$$\begin{bmatrix} 1 & 0 & 1 & 2 \\ 0 & 2 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 2 & 0 & 1 & 1 \end{bmatrix}$$

Filter K (edge detector):

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

Step 1: Position filter at top-left:

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 2 & 1 \\ 1 & 1 & 0 \end{bmatrix} \circ \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\begin{aligned} &= (1 \times -1) + (0 \times 0) + (1 \times 1) + (0 \times -1) + (2 \times 0) + (1 \times 1) + (1 \times -1) + (1 \times 0) + (0 \times 1) \\ &= -1 + 0 + 1 + 0 + 0 + 1 - 1 + 0 + 0 = 0 \end{aligned}$$

Step 2: Slide right and continue. Final output (2×2):

$$\begin{bmatrix} 0 & 3 \\ 2 & -1 \end{bmatrix}$$

3.4 Output Size Formula

Given:

- Input size: $W \times H$
- Filter size: $F \times F$
- Padding: P (extra pixels around image)
- Stride: S (how many pixels to jump)

$$\text{Output Width} = \frac{W - F + 2P}{S} + 1$$

$$\text{Output Height} = \frac{H - F + 2P}{S} + 1$$

Example: Input 6×6 , Filter 3×3 , Padding 0, Stride 1:

$$\frac{6 - 3 + 0}{1} + 1 = 4 \Rightarrow 4 \times 4 \text{ output}$$

Example: Input 6×6 , Filter 3×3 , Padding 0, Stride 2:

$$\frac{6 - 3 + 0}{2} + 1 = 2.5 \rightarrow 2 \Rightarrow 2 \times 2 \text{ output}$$

4 Pooling Layer: The Downsample

4.1 Why Pooling?

1. Reduce size (faster computation)
2. Make features translation invariant
3. Prevent overfitting

4.2 Types of Pooling

4.2.1 Max Pooling (Most Common)

Takes maximum value in each window.

Example: 4×4 input, 2×2 window, stride 2:

$$\begin{bmatrix} 1 & 3 & 2 & 9 \\ 4 & 6 & 1 & 5 \\ 7 & 2 & 8 & 3 \\ 0 & 4 & 3 & 6 \end{bmatrix} \rightarrow \begin{bmatrix} 6 & 9 \\ 7 & 8 \end{bmatrix}$$

Explanation:

- Window 1 (top-left): $\max(1,3,4,6) = 6$
- Window 2 (top-right): $\max(2,9,1,5) = 9$
- Window 3 (bottom-left): $\max(7,2,0,4) = 7$
- Window 4 (bottom-right): $\max(8,3,3,6) = 8$

4.2.2 Average Pooling

Takes average value in each window.

Same example:

$$\begin{bmatrix} 1 & 3 & 2 & 9 \\ 4 & 6 & 1 & 5 \\ 7 & 2 & 8 & 3 \\ 0 & 4 & 3 & 6 \end{bmatrix} \rightarrow \begin{bmatrix} (1+3+4+6)/4 & (2+9+1+5)/4 \\ (7+2+0+4)/4 & (8+3+3+6)/4 \end{bmatrix} = \begin{bmatrix} 3.5 & 4.25 \\ 3.25 & 5.0 \end{bmatrix}$$

4.2.3 Min Pooling (Rare)

Takes minimum value in each window.

5 Complete CNN Example: Step-by-Step

5.1 Input Image

RGB image: $32 \times 32 \times 3$ (32 pixels tall, 32 wide, 3 color channels)

5.2 Layer 1: Convolution

- 32 filters of size $3 \times 3 \times 3$
- Stride = 1, Padding = 1
- Output size calculation:

$$\frac{32 - 3 + 2 \times 1}{1} + 1 = 32 \Rightarrow 32 \times 32 \times 32$$

(32 feature maps, each 32×32)

5.3 Layer 2: Max Pooling

- Window: 2×2
- Stride = 2
- Output size:

$$\frac{32}{2} = 16 \Rightarrow 16 \times 16 \times 32$$

5.4 Layer 3: Convolution

- 64 filters of size $3 \times 3 \times 32$
- Stride = 1, Padding = 1
- Output: $16 \times 16 \times 64$

5.5 Layer 4: Max Pooling

- Window: 2×2 , Stride = 2
- Output: $8 \times 8 \times 64$

5.6 Layer 5: Flatten

Convert 3D tensor to 1D vector:

$$8 \times 8 \times 64 = 4,096 \text{ neurons}$$

5.7 Layer 6: Fully Connected

- $4,096 \rightarrow 128$ neurons
- Uses ReLU activation

5.8 Layer 7: Output Layer

- $128 \rightarrow 10$ neurons (for 10 classes)
- Uses Softmax activation

6 Activation Functions

6.1 ReLU (Rectified Linear Unit)

$$f(x) = \max(0, x)$$

- Pros: Simple, fast, prevents vanishing gradient
- Cons: "Dying ReLU" problem

6.2 Softmax (for output layer)

Converts scores to probabilities:

$$P(\text{class } i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

Where C = number of classes.

Example: Scores = [2.0, 1.0, 0.1]

$$\text{Sum} = e^{2.0} + e^{1.0} + e^{0.1} = 7.39 + 2.72 + 1.11 = 11.22$$

$$P(\text{class 1}) = \frac{7.39}{11.22} = 0.66$$

$$P(\text{class 2}) = \frac{2.72}{11.22} = 0.24$$

$$P(\text{class 3}) = \frac{1.11}{11.22} = 0.10$$

7 Training a CNN

7.1 Loss Function: Cross-Entropy

For classification:

$$L = - \sum_{i=1}^C y_i \log(\hat{y}_i)$$

Where:

- y_i = true label (1 for correct class, 0 others)
- \hat{y}_i = predicted probability

7.2 Backpropagation in CNNs

The chain rule applied through:

1. Output layer gradients
2. Fully connected layers
3. Pooling layers (pass gradients to max position)
4. Convolution layers (update filter weights)

7.3 Weight Update

Using gradient descent:

$$W_{\text{new}} = W_{\text{old}} - \eta \frac{\partial L}{\partial W}$$

Where η = learning rate.

8 Advanced CNN Concepts

8.1 1×1 Convolutions

- Used for dimensionality reduction
- Example: $32 \times 32 \times 256 \rightarrow 32 \times 32 \times 64$ using 64 filters of size $1 \times 1 \times 256$
- Reduces computation cost

8.2 Dilated Convolutions

- Increases receptive field without pooling
- Has "holes" in filter
- Used in segmentation tasks

8.3 Batch Normalization

Normalizes layer inputs:

$$\hat{x} = \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

Reduces internal covariate shift, speeds training.

9 Practical Example: Classifying Cats vs Dogs

1. **Input:** $128 \times 128 \times 3$ (RGB image)
2. **Conv1:** 32 filters, $3 \times 3 \rightarrow 128 \times 128 \times 32$
3. **Pool1:** 2×2 max pool $\rightarrow 64 \times 64 \times 32$
4. **Conv2:** 64 filters, $3 \times 3 \rightarrow 64 \times 64 \times 64$
5. **Pool2:** 2×2 max pool $\rightarrow 32 \times 32 \times 64$
6. **Conv3:** 128 filters, $3 \times 3 \rightarrow 32 \times 32 \times 128$
7. **Pool3:** 2×2 max pool $\rightarrow 16 \times 16 \times 128$
8. **Flatten:** $16 \times 16 \times 128 = 32,768$ neurons
9. **FC1:** $32,768 \rightarrow 512$ (ReLU)
10. **FC2:** $512 \rightarrow 2$ (Softmax)

10 Common CNN Architectures

10.1 LeNet-5 (1998)

- First successful CNN
- For digit recognition
- 7 layers total

10.2 AlexNet (2012)

- Won ImageNet competition
- 8 layers, ReLU activation
- Dropout for regularization

10.3 VGG-16 (2014)

- Simple: only 3×3 convolutions
- 16 layers deep
- 138 million parameters

10.4 ResNet (2015)

- Residual connections
- Solves vanishing gradient in deep networks
- Up to 152 layers

11 Mathematics Deep Dive

11.1 Convolution as Matrix Multiplication

A convolution can be expressed as matrix multiplication using the **im2col** operation:

Example: 3×3 input, 2×2 filter:

$$\begin{bmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{bmatrix} \rightarrow \begin{bmatrix} x_{11} & x_{12} & x_{21} & x_{22} \\ x_{12} & x_{13} & x_{22} & x_{23} \\ x_{21} & x_{22} & x_{31} & x_{32} \\ x_{22} & x_{23} & x_{32} & x_{33} \end{bmatrix}$$

Filter weights as column vector:

$$\begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \rightarrow \begin{bmatrix} w_{11} \\ w_{12} \\ w_{21} \\ w_{22} \end{bmatrix}$$

Convolution = Matrix multiplication:

$$\text{Output} = \text{im2col}(X) \times W$$

11.2 Gradient Calculation for Convolution

For a convolution layer:

$$\frac{\partial L}{\partial K} = \text{convolution} \left(\frac{\partial L}{\partial O}, I \right)$$

Where:

- K = filter weights
- O = output
- I = input

12 Implementation Tips

1. **Always normalize inputs:** Scale pixels to $[0,1]$ or $[-1,1]$
2. **Use data augmentation:** Rotation, flipping, zooming
3. **Start with small filters:** 3×3 is standard
4. **Increase filters as network deepens:** $32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
5. **Use dropout:** Prevents overfitting
6. **Monitor training/validation loss:** Stop when validation loss increases

13 Summary

- **CNN** = Feature extraction + Classification
- **Convolution** = Pattern detection with shared weights
- **Pooling** = Downsampling for translation invariance
- **Training** = Minimize loss via backpropagation
- **Key advantage:** Parameter sharing → efficient for images

14 Exercises

14.1 Exercise 1: Size Calculation

Calculate output sizes:

1. Input: $28 \times 28 \times 1$, Filter: 5×5 , Stride: 1, Padding: 0
2. Input: $100 \times 100 \times 3$, Filter: 3×3 , Stride: 2, Padding: 1
3. Input: $14 \times 14 \times 64$ after 2×2 max pooling with stride 2

14.2 Exercise 2: Convolution Calculation

Perform convolution:

$$I = \begin{bmatrix} 2 & 4 & 1 \\ 3 & 1 & 5 \\ 0 & 2 & 1 \end{bmatrix}, \quad K = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

Stride = 1, no padding.

14.3 Exercise 3: Pooling

Apply max pooling (2×2 , stride 2):

$$\begin{bmatrix} 3 & 8 & 2 & 4 \\ 1 & 6 & 5 & 3 \\ 7 & 2 & 9 & 1 \\ 4 & 5 & 3 & 6 \end{bmatrix}$$

Answers

Exercise 1

1. $\frac{28-5}{1} + 1 = 24 \rightarrow 24 \times 24$
2. $\frac{100-3+2}{2} + 1 = 50 \rightarrow 50 \times 50 \times ?$ (depends on filters)
3. $\frac{14}{2} = 7 \rightarrow 7 \times 7 \times 64$

Exercise 2

$$\begin{bmatrix} (2 \times 0) + (4 \times 1) + (3 \times -1) + (1 \times 0) & (4 \times 0) + (1 \times 1) + (1 \times -1) + (5 \times 0) \\ (3 \times 0) + (1 \times 1) + (0 \times -1) + (2 \times 0) & (1 \times 0) + (5 \times 1) + (2 \times -1) + (1 \times 0) \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 1 & 3 \end{bmatrix}$$

Exercise 3

$$\begin{bmatrix} \max(3, 8, 1, 6) & \max(2, 4, 5, 3) \\ \max(7, 2, 4, 5) & \max(9, 1, 3, 6) \end{bmatrix} = \begin{bmatrix} 8 & 5 \\ 7 & 9 \end{bmatrix}$$