

# Notes

## Perceptron vs ANN vs CNN on MNIST Dataset

Ankit Kumar

February 4, 2026

## Contents

<b>1</b>	<b>Introduction and Theoretical Foundations</b>	<b>3</b>
1.1	Evolution of Neural Networks . . . . .	3
1.2	Mathematical Foundations . . . . .	3
1.2.1	Perceptron Mathematical Model . . . . .	3
1.2.2	Multi-layer Neural Networks . . . . .	3
1.2.3	Backpropagation Algorithm . . . . .	3
1.3	Activation Functions Comparison . . . . .	4
<b>2</b>	<b>Implementation Setup and Environment</b>	<b>4</b>
2.1	Complete Import Statements . . . . .	4
2.2	Data Loading and Exploration . . . . .	5
2.3	Data Preprocessing Steps . . . . .	5
<b>3</b>	<b>Perceptron Implementation</b>	<b>7</b>
3.1	Perceptron Architecture Design . . . . .	7
3.2	Model Compilation . . . . .	7
3.3	Model Training . . . . .	8
3.4	Training Output Analysis . . . . .	8
3.4.1	Key Observations . . . . .	9
3.5	Performance Evaluation . . . . .	9
<b>4</b>	<b>Artificial Neural Network (ANN) Implementation</b>	<b>10</b>
4.1	ANN Architecture Design . . . . .	10
4.2	ANN Compilation with Adam Optimizer . . . . .	11
4.3	ANN Training . . . . .	11
4.4	ANN Training Output Analysis . . . . .	12
4.4.1	Key Observations . . . . .	12
4.5	ANN Evaluation . . . . .	12
<b>5</b>	<b>Convolutional Neural Network (CNN) Implementation</b>	<b>12</b>
5.1	Data Preparation for CNN . . . . .	12
5.2	CNN Architecture Design . . . . .	13
5.3	Convolution Operation Details . . . . .	14
5.3.1	Convolution Mathematics . . . . .	14

5.3.2	Output Size Calculation . . . . .	15
5.3.3	Pooling Operation . . . . .	15
5.4	CNN Compilation and Training . . . . .	15
5.5	CNN Training Output Analysis . . . . .	15
5.5.1	Key Observations . . . . .	16
5.6	CNN Evaluation . . . . .	16
<b>6</b>	<b>Visualization and Analysis Functions</b>	<b>16</b>
<b>7</b>	<b>Comprehensive Analysis and Comparison</b>	<b>18</b>
7.1	Parameter Count Analysis . . . . .	18
7.2	Performance Metrics Summary . . . . .	18
7.3	Training Time Analysis . . . . .	18
7.4	Error Analysis by Digit . . . . .	19
<b>8</b>	<b>Key Insights and Learnings</b>	<b>19</b>
8.1	Architectural Insights . . . . .	19
8.2	Practical Recommendations . . . . .	20
8.3	Limitations and Future Improvements . . . . .	20
<b>9</b>	<b>Conclusion</b>	<b>21</b>
9.1	Key Takeaways . . . . .	21
9.2	Final Performance Ranking . . . . .	21
9.3	Educational Value . . . . .	21
9.4	Code Availability . . . . .	21

# 1 Introduction and Theoretical Foundations

## 1.1 Evolution of Neural Networks

The journey from perceptrons to deep learning represents one of the most significant developments in artificial intelligence:

- **1957 - Perceptron:** Single-layer linear classifier by Frank Rosenblatt
- **1960s-1970s:** AI winter due to perceptron limitations (XOR problem)
- **1980s:** Backpropagation algorithm enables multi-layer networks
- **1990s:** Convolutional Neural Networks for image processing
- **2000s-present:** Deep learning revolution with improved hardware and algorithms

## 1.2 Mathematical Foundations

### 1.2.1 Perceptron Mathematical Model

For input vector  $\mathbf{x} = [x_1, x_2, \dots, x_n]$ , weights  $\mathbf{w} = [w_1, w_2, \dots, w_n]$ , and bias  $b$ :

$$y = f \left( \sum_{i=1}^n w_i x_i + b \right) \quad (1)$$

where  $f$  is typically a step function:  $f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$

### 1.2.2 Multi-layer Neural Networks

For a network with  $L$  layers:

$$\mathbf{z}^{(l)} = \mathbf{W}^{(l)} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \quad (2)$$

$$\mathbf{a}^{(l)} = f(\mathbf{z}^{(l)}) \quad (3)$$

where  $\mathbf{a}^{(0)} = \mathbf{x}$  is the input.

### 1.2.3 Backpropagation Algorithm

The weight update rule using gradient descent:

$$\mathbf{W}^{(l)} := \mathbf{W}^{(l)} - \alpha \frac{\partial J}{\partial \mathbf{W}^{(l)}} \quad (4)$$

where  $\alpha$  is the learning rate and  $J$  is the cost function.

Function	Formula	Range	Use Case
Sigmoid	$\sigma(x) = \frac{1}{1+e^{-x}}$	(0,1)	Output layer, probability
Tanh	$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	(-1,1)	Hidden layers
ReLU	$\max(0, x)$	[0,)	Hidden layers (prevents vanishing gradient)
Softmax	$\frac{e^{x_i}}{\sum_j e^{x_j}}$	(0,1), sums to 1	Multi-class output

Table 1: Common Activation Functions

### 1.3 Activation Functions Comparison

## 2 Implementation Setup and Environment

### 2.1 Complete Import Statements

```

1 # =====
2 # BASIC IMPORTS FOR DATA MANIPULATION
3 # =====
4 import numpy as np # Numerical computations
5 import pandas as pd # Data handling
6 import seaborn as sns # Statistical visualization
7 import matplotlib.pyplot as plt # Plotting
8
9 # Suppress warnings for cleaner output
10 import warnings
11 warnings.filterwarnings('ignore')
12
13 # =====
14 # SCIKIT-LEARN IMPORTS
15 # =====
16 from sklearn.preprocessing import LabelEncoder, StandardScaler
17 # LabelEncoder: Converts categorical labels to numerical
18 # StandardScaler: Standardizes features by removing mean and scaling
19
20 from sklearn.model_selection import train_test_split
21 # Splits data into training and testing sets
22
23 from sklearn.linear_model import Perceptron
24 # Classical perceptron implementation from scikit-learn
25
26 from sklearn.metrics import accuracy_score, confusion_matrix,
27     classification_report
28 # accuracy_score: Calculates accuracy of predictions
29 # confusion_matrix: Shows correct/incorrect classifications
30 # classification_report: Detailed performance metrics
31
32 # =====
33 # TENSORFLOW/KERAS IMPORTS
34 # =====
35 from tensorflow.keras.models import Sequential
36 # Sequential model: Linear stack of layers
37
38 from tensorflow.keras.layers import Dense, Conv2D, Flatten,
39     MaxPooling2D, Dropout
# Dense: Fully connected layer
# Conv2D: 2D convolutional layer

```

```

40 # Flatten: Converts multi-dimensional input to 1D
41 # MaxPooling2D: Max pooling operation
42 # Dropout: Regularization layer that randomly sets inputs to zero
43
44 from tensorflow.keras.utils import to_categorical
45 # Converts class vectors to binary class matrix (one-hot encoding)

```

Listing 1: Complete Python Imports

## 2.2 Data Loading and Exploration

```

1 # Load MNIST dataset from CSV files
2 df = pd.read_csv('/content/sample_data/mnist_train_small.csv')
3 df_test = pd.read_csv('/content/sample_data/mnist_test.csv')
4
5 # Display dataset shapes
6 print("Training Data Shape:", df.shape) # Expected: (19999, 785)
7 print("Test Data Shape:", df_test.shape) # Expected: (10000, 785)
8
9 # Explanation of the dataset structure:
10 # - 785 columns: 1 label column + 784 pixel columns (28 28 image)
11 # - Each pixel value ranges from 0 to 255 (grayscale intensity)
12 # - Labels: 0-9 representing handwritten digits
13
14 # Display column names (first 10)
15 print("\nFirst 10 column names:")
16 print(df.columns[:10])
17 # Note: The CSV doesn't have proper column headers,
18 # so pandas assigns default numeric names
19
20 # Extract features and labels
21 # Since columns don't have proper names, we use .iloc for positional
22 # indexing
23 # Column 0: Label (digit 0-9)
24 # Columns 1-784: Pixel values
25 x_train = df.iloc[:, 1:].values # All rows, columns 1 to end (pixels)
26 y_train = df.iloc[:, 0].values # All rows, column 0 (labels)
27
28 x_test = df_test.iloc[:, 1:].values
29 y_test = df_test.iloc[:, 0].values
30
31 print(f"\nTraining data shape: {x_train.shape}") # (19999, 784)
32 print(f"Training labels shape: {y_train.shape}") # (19999,)
33 print(f"Test data shape: {x_test.shape}") # (10000, 784)
34 print(f"Test labels shape: {y_test.shape}") # (10000,)

```

Listing 2: Loading and Exploring MNIST Dataset

## 2.3 Data Preprocessing Steps

```

1 # =====
2 # STEP 1: DATA NORMALIZATION
3 # =====
4 # Convert to float32 and normalize pixel values from [0,255] to [0,1]
5 # Why normalize?

```

```

6 # 1. Prevents numerical instability during training
7 # 2. Helps gradient descent converge faster
8 # 3. Ensures consistent scale across features
9 x_train = x_train.astype("float32") / 255.0
10 x_test = x_test.astype("float32") / 255.0
11
12 print("Normalized pixel value range:")
13 print(f"Min: {x_train.min()}, Max: {x_train.max()}")
14 print(f"Mean: {x_train.mean():.4f}, Std: {x_train.std():.4f}")
15
16 # =====
17 # STEP 2: RESHAPING FOR VISUALIZATION
18 # =====
19 # Reshape from (samples, 784) to (samples, 28, 28) for 2D
20 # representation
21 # This doesn't change the data, just reorganizes it for easier handling
22 x_train_img = x_train.reshape(-1, 28, 28) # -1 means infer this
23 # dimension
24 x_test_img = x_test.reshape(-1, 28, 28)
25
26 print(f"\nReshaped training data: {x_train_img.shape}") # (19999, 28,
27 28)
28 print(f"Reshaped test data: {x_test_img.shape}") # (10000, 28,
29 28)
30
31 # Verify reshaping by looking at first image
32 print(f"\nFirst image shape: {x_train_img[0].shape}")
33 print(f"First image min pixel: {x_train_img[0].min()}")
34 print(f"First image max pixel: {x_train_img[0].max()}")
35
36 # =====
37 # STEP 3: LABEL ENCODING (ONE-HOT ENCODING)
38 # =====
39 # Convert integer labels (0-9) to one-hot encoded vectors
40 # Example: label 3 becomes [0, 0, 0, 1, 0, 0, 0, 0, 0, 0]
41 # Why one-hot encoding?
42 # 1. Required for categorical cross-entropy loss
43 # 2. Each class gets its own output neuron
44 # 3. Better for multi-class classification
45 y_train_cat = to_categorical(y_train, 10) # 10 classes (digits 0-9)
46 y_test_cat = to_categorical(y_test, 10)
47
48 print(f"\nOriginal label 0: {y_train[0]}")
49 print(f"One-hot encoded label 0: {y_train_cat[0]}")
50 print(f"One-hot encoded shape: {y_train_cat.shape}") # (19999, 10)
51
52 # Verify one-hot encoding
53 print(f"\nSum of one-hot vector (should be 1): {y_train_cat[0].sum()}")
54 print(f"Number of unique labels: {len(np.unique(y_train))}")

```

Listing 3: Data Preprocessing Pipeline

### 3 Perceptron Implementation

#### 3.1 Perceptron Architecture Design

The perceptron implemented here is a single-layer neural network with:

- **Input:** 784 neurons (flattened  $28 \times 28$  image)
- **Output:** 10 neurons with softmax activation (one per digit class)
- **Total Parameters:** 7,850 ( $784 \times 10$  weights + 10 biases)
- **Activation:** Softmax for multi-class probability distribution

```
1 # =====
2 # PERCEPTRON MODEL DEFINITION
3 # =====
4 # Using Keras Sequential API for simple linear stack of layers
5 Perceptron = Sequential([
6     # Layer 1: Flatten layer
7     # Converts 28 28 input image to 1D vector of 784 elements
8     # Required because Dense layers expect 1D input
9     Flatten(input_shape=(28,28)),    # Output shape: (batch_size, 784)
10
11    # Layer 2: Dense (fully connected) output layer
12    # 10 neurons for 10 digit classes (0-9)
13    # softmax activation converts outputs to probability distribution
14    # softmax ensures all outputs sum to 1
15    Dense(10, activation='softmax')  # Output shape: (batch_size, 10)
16])
17
18 # Display model architecture summary
19 print("*"*50)
20 print("PERCEPTRON MODEL SUMMARY")
21 print("*"*50)
22 Perceptron.summary()
23
24 # Expected output:
25 # Total params: 7,850
26 # Trainable params: 7,850
27 # Non-trainable params: 0
```

Listing 4: Perceptron Model Definition

#### 3.2 Model Compilation

```
1 # =====
2 # MODEL COMPIILATION
3 # =====
4 # Compile configures the model for training
5 Perceptron.compile(
6     optimizer='sgd',  # Stochastic Gradient Descent
7     # SGD updates weights after each training example
8     # Simple but effective for linear models
9 )
```

```

10 loss='categorical_crossentropy', # Loss function for multi-class
11 # Measures difference between predicted and true distributions
12 # Formula: -      y_true * log(y_pred)
13
14 metrics=['accuracy'] # Evaluation metric to monitor
15 # Accuracy = (correct predictions) / (total predictions)
16 )
17
18 # Alternative optimizer configurations:
19 # optimizer='sgd' - basic stochastic gradient descent
20 # optimizer=tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9)
21 # optimizer='adam' - adaptive moment estimation (better for complex
22 # models)

```

Listing 5: Perceptron Model Compilation

### 3.3 Model Training

```

1 # =====
2 # MODEL TRAINING
3 # =====
4 # Fit method trains the model on training data
5 history_precp = Perceptron.fit(
6     x_train_img,          # Training features (images)
7     y_train_cat,         # Training labels (one-hot encoded)
8
9     epochs=10,           # Number of complete passes through training data
10    # Each epoch = one forward pass + one backward pass of ALL training
11    # samples
12    # More epochs = more learning, but risk of overfitting
13
14    batch_size=32,        # Number of samples per gradient update
15    # Smaller batches: more frequent updates, noisier gradient
16    # Larger batches: smoother gradient, more memory required
17
18    validation_data=(x_test_img, y_test_cat), # Validation/test data
19    # Used to evaluate loss and metrics after each epoch
20    # Not used for training, only for monitoring
21
22    verbose=1            # Verbosity mode (1 = progress bar)
)
23
24 # Training process explanation:
25 # 1. Forward pass: Compute predictions
26 # 2. Compute loss: Compare predictions with true labels
27 # 3. Backward pass: Compute gradients using backpropagation
28 # 4. Update weights: Adjust weights using optimizer
29 # 5. Repeat for all batches in epoch
30 # 6. Repeat for all epochs

```

Listing 6: Perceptron Training Process

### 3.4 Training Output Analysis

The perceptron training shows typical learning behavior:

```

Epoch 1/10: accuracy: 0.5885 loss: 1.5343 val_accuracy: 0.8497
Epoch 2/10: accuracy: 0.8489 loss: 0.6749 val_accuracy: 0.8720
Epoch 3/10: accuracy: 0.8715 loss: 0.5393 val_accuracy: 0.8818
Epoch 4/10: accuracy: 0.8784 loss: 0.4835 val_accuracy: 0.8903
Epoch 5/10: accuracy: 0.8868 loss: 0.4465 val_accuracy: 0.8932
Epoch 6/10: accuracy: 0.8880 loss: 0.4243 val_accuracy: 0.8964
Epoch 7/10: accuracy: 0.8884 loss: 0.4136 val_accuracy: 0.8986
Epoch 8/10: accuracy: 0.8948 loss: 0.3957 val_accuracy: 0.9005
Epoch 9/10: accuracy: 0.8949 loss: 0.3912 val_accuracy: 0.9025
Epoch 10/10: accuracy: 0.8976 loss: 0.3759 val_accuracy: 0.9032

```

### 3.4.1 Key Observations

- **Rapid initial learning:** Accuracy jumps from 58.85% to 84.89% in first epoch
- **Gradual improvement:** Slower learning in later epochs
- **Good generalization:** Validation accuracy (90.32%) close to training accuracy (89.76%)
- **Stable convergence:** Loss steadily decreases without oscillations

## 3.5 Performance Evaluation

```

1 # =====
2 # PERFORMANCE EVALUATION
3 # =====
4 # Evaluate model on test data
5 test_loss, test_accuracy = Perceptron.evaluate(x_test_img, y_test_cat,
6     verbose=0)
7 print("\n" + "="*50)
8 print("PERCEPTRON FINAL EVALUATION")
9 print("="*50)
10 print(f"Test Loss: {test_loss:.4f}")
11 print(f"Test Accuracy: {test_accuracy:.4f} ({test_accuracy*100:.2f}%)")
12
13 # Alternative evaluation methods:
14 # 1. Get predictions
15 predictions = Perceptron.predict(x_test_img)
16 predicted_classes = np.argmax(predictions, axis=1)
17 true_classes = np.argmax(y_test_cat, axis=1)
18
19 # 2. Calculate accuracy manually
20 correct_predictions = np.sum(predicted_classes == true_classes)
21 total_predictions = len(true_classes)
22 manual_accuracy = correct_predictions / total_predictions
23 print(f"\nManual Accuracy Calculation:")
24 print(f"Correct: {correct_predictions}/{total_predictions}")
25 print(f"Accuracy: {manual_accuracy:.4f}")
26
27 # 3. Confusion Matrix
28 from sklearn.metrics import confusion_matrix
29 cm = confusion_matrix(true_classes, predicted_classes)
30 print(f"\nConfusion Matrix Shape: {cm.shape}")

```

```

30 print("Confusion Matrix (first 5x5):")
31 print(cm[:5, :5])

```

Listing 7: Perceptron Evaluation

## 4 Artificial Neural Network (ANN) Implementation

### 4.1 ANN Architecture Design

The ANN extends the perceptron with hidden layers:

- **Input:** 784 neurons (flattened image)
- **Hidden Layer 1:** 128 neurons with ReLU activation
- **Hidden Layer 2:** 64 neurons with ReLU activation
- **Output:** 10 neurons with softmax activation
- **Total Parameters:** 109,386

```

1 # =====
2 # ARTIFICIAL NEURAL NETWORK (ANN) DEFINITION
3 # =====
4 ann = Sequential([
5     # Layer 1: Flatten input
6     Flatten(input_shape=(28,28)),    # (batch_size, 784)
7
8     # Layer 2: First hidden layer
9     Dense(128, activation='relu'),    # (batch_size, 128)
10    # 128 neurons with ReLU activation
11    # ReLU: f(x) = max(0, x)
12    # Advantages: Computationally efficient, mitigates vanishing
13    # gradient
14
15    # Layer 3: Second hidden layer
16    Dense(64, activation='relu'),    # (batch_size, 64)
17    # 64 neurons with ReLU activation
18    # Creates deeper feature representations
19
20    # Layer 4: Output layer
21    Dense(10, activation='softmax') # (batch_size, 10)
22    # 10 neurons for digit classification
23])
24
25 # Calculate total parameters:
26 # Layer 1 (Flatten): 0 parameters (only reshapes)
27 # Layer 2 (Dense 128): (784 128 ) weights + 128 biases = 100,480
28 # Layer 3 (Dense 64): (128 64 ) weights + 64 biases = 8,256
29 # Layer 4 (Dense 10): (64 10 ) weights + 10 biases = 650
30 # Total: 100,480 + 8,256 + 650 = 109,386
31
32 print("*"*50)
33 print("ARTIFICIAL NEURAL NETWORK SUMMARY")
34 print("*"*50)

```

```
34 ann.summary()
```

Listing 8: ANN Model Definition

## 4.2 ANN Compilation with Adam Optimizer

```
1 # =====
2 # ANN COMPILATION
3 # =====
4 ann.compile(
5     optimizer='adam', # Adaptive Moment Estimation
6     # Adam combines benefits of two extensions of SGD:
7     # 1. Adaptive Gradient Algorithm (AdaGrad): Maintains per-parameter
8     # learning rates
9     # 2. Root Mean Square Propagation (RMSPProp): Maintains moving
10    # average of squared gradients
11    # Advantages: Faster convergence, less sensitive to learning rate
12
13    loss='categorical_crossentropy',
14    metrics=['accuracy']
15)
16
17 # Adam optimizer parameters (default values):
18 # learning_rate=0.001
19 # beta_1=0.9 (exponential decay rate for first moment estimates)
20 # beta_2=0.999 (exponential decay rate for second moment estimates)
21 # epsilon=1e-07 (numerical stability term)
22
23 # Custom Adam optimizer:
# from tensorflow.keras.optimizers import Adam
# optimizer=Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999)
```

Listing 9: ANN Compilation with Adam Optimizer

## 4.3 ANN Training

```
1 # =====
2 # ANN TRAINING
3 # =====
4 history_ann = ann.fit(
5     x_train_img,
6     y_train_cat,
7     epochs=10,
8     batch_size=32,
9     validation_data=(x_test_img, y_test_cat),
10    verbose=1
11)
12
13 # Training behavior expectations:
14 # 1. Faster convergence than perceptron due to Adam optimizer
15 # 2. Higher capacity to learn complex patterns (more parameters)
16 # 3. Risk of overfitting due to increased model complexity
17 # 4. Better feature learning through hidden layers
```

Listing 10: ANN Training Process

## 4.4 ANN Training Output Analysis

```
Epoch 1/10: accuracy: 0.8126 loss: 0.6513 val_accuracy: 0.9403
Epoch 2/10: accuracy: 0.9525 loss: 0.1713 val_accuracy: 0.9509
Epoch 3/10: accuracy: 0.9674 loss: 0.1141 val_accuracy: 0.9586
Epoch 4/10: accuracy: 0.9753 loss: 0.0806 val_accuracy: 0.9607
Epoch 5/10: accuracy: 0.9796 loss: 0.0632 val_accuracy: 0.9637
Epoch 6/10: accuracy: 0.9867 loss: 0.0446 val_accuracy: 0.9689
Epoch 7/10: accuracy: 0.9900 loss: 0.0334 val_accuracy: 0.9672
Epoch 8/10: accuracy: 0.9941 loss: 0.0219 val_accuracy: 0.9667
Epoch 9/10: accuracy: 0.9952 loss: 0.0196 val_accuracy: 0.9644
Epoch 10/10: accuracy: 0.9944 loss: 0.0164 val_accuracy: 0.9598
```

### 4.4.1 Key Observations

- **Faster convergence:** Reaches 94.03% validation accuracy in just 1 epoch
- **Higher training accuracy:** 99.44% vs perceptron's 89.76%
- **Better generalization:** 95.98% validation accuracy vs 90.32%
- **Overfitting signs:** Training accuracy (99.44%) < Validation accuracy (95.98%)
- **Lower loss:** 0.1620 vs perceptron's 0.3613

## 4.5 ANN Evaluation

```
1 # =====
2 # ANN EVALUATION
3 # =====
4 acc_ann = ann.evaluate(x_test_img, y_test_cat, verbose=0)[1]
5 print("\n" + "="*50)
6 print("ARTIFICIAL NEURAL NETWORK FINAL EVALUATION")
7 print("="*50)
8 print(f"Test Accuracy: {acc_ann:.4f} ({acc_ann*100:.2f}%)")
9
10 # Performance comparison with perceptron
11 print(f"\nPerformance Improvement over Perceptron:")
12 print(f"Accuracy increase: {((acc_ann - acc_precp)*100:.2f}%)")
13 print(f"Relative improvement: {((acc_ann/acc_precp)-1)*100:.2f}%)")
```

Listing 11: ANN Performance Evaluation

## 5 Convolutional Neural Network (CNN) Implementation

### 5.1 Data Preparation for CNN

CNNs require 4D input tensors: (samples, height, width, channels)

```

1 # =====
2 # DATA PREPARATION FOR CNN
3 # =====
4 # CNN expects 4D input: (batch_size, height, width, channels)
5 # MNIST images are grayscale, so channels=1
6 X_train_cnn = x_train.reshape(-1, 28, 28, 1) # Shape: (19999, 28, 28,
7 1)
8 X_test_cnn = x_test.reshape(-1, 28, 28, 1) # Shape: (10000, 28, 28,
9 1)
10
11 print("CNN Input Shapes:")
12 print(f"Training: {X_train_cnn.shape}")
13 print(f"Test: {X_test_cnn.shape}")
14
15 # Explanation of dimensions:
16 # Dimension 0: Number of samples (19999 for training)
17 # Dimension 1: Height (28 pixels)
18 # Dimension 2: Width (28 pixels)
19 # Dimension 3: Channels (1 for grayscale, 3 for RGB)
20
21 # For color images, you would reshape to (samples, height, width, 3)
22 # Example: X_train_cnn = x_train.reshape(-1, 28, 28, 3)

```

Listing 12: Data Reshaping for CNN

## 5.2 CNN Architecture Design

```

1 # =====
2 # CONVOLUTIONAL NEURAL NETWORK DEFINITION
3 # =====
4 cnn = Sequential([
5     # ===== CONVOLUTIONAL BLOCK 1 =====
6     # Layer 1: First convolutional layer
7     Conv2D(32, kernel_size=(3,3), activation='relu', input_shape
8         =(28,28,1)),
9     # 32 filters, each 3 3
10    # Input: (batch_size, 28, 28, 1)
11    # Output: (batch_size, 26, 26, 32) [28-3+1=26, no padding]
12    # Parameters: (3      3132      ) + 32 = 320
13
14    # Layer 2: First pooling layer
15    MaxPooling2D(pool_size=(2,2)),
16    # Max pooling with 2 2 window, stride=2
17    # Input: (batch_size, 26, 26, 32)
18    # Output: (batch_size, 13, 13, 32) [26/2=13]
19    # Reduces spatial dimensions by half
20
21    # ===== CONVOLUTIONAL BLOCK 2 =====
22    # Layer 3: Second convolutional layer
23    Conv2D(64, kernel_size=(3,3), activation='relu'),
24    # 64 filters, each 3 3
25    # Input: (batch_size, 13, 13, 32)
26    # Output: (batch_size, 11, 11, 64) [13-3+1=11]
27    # Parameters: (3      33264      ) + 64 = 18,496
28
29    # Layer 4: Second pooling layer

```

```

29     MaxPooling2D(pool_size=(2,2)),
30     # Input: (batch_size, 11, 11, 64)
31     # Output: (batch_size, 5, 5, 64) [11/2=5.5      5]
32     # Note: 11/2 = 5.5, integer division gives 5
33
34     # ===== FLATTEN AND DENSE LAYERS =====
35     # Layer 5: Flatten layer
36     Flatten(),
37     # Converts 3D feature maps to 1D vector
38     # Input: (batch_size, 5, 5, 64)
39     # Output: (batch_size, 1600) [5 564 =1600]
40
41     # Layer 6: Fully connected layer
42     Dense(128, activation='relu'),
43     # 128 neurons with ReLU activation
44     # Parameters: (1600 128 ) + 128 = 204,928
45
46     # Layer 7: Dropout for regularization
47     Dropout(0.5),
48     # Randomly sets 50% of inputs to zero during training
49     # Prevents overfitting by preventing co-adaptation of neurons
50     # No parameters
51
52     # Layer 8: Output layer
53     Dense(10, activation='softmax')
54     # 10 neurons for digit classification
55     # Parameters: (128 10 ) + 10 = 1,290
56 )
57
58 # Total parameters calculation:
59 # Conv2D 1: 320
60 # Conv2D 2: 18,496
61 # Dense 1: 204,928
62 # Dense 2: 1,290
63 # Total: 225,034
64
65 print("*"*50)
66 print("CONVOLUTIONAL NEURAL NETWORK SUMMARY")
67 print("*"*50)
68 cnn.summary()

```

Listing 13: CNN Model Definition

## 5.3 Convolution Operation Details

### 5.3.1 Convolution Mathematics

For a single filter  $F$  of size  $k \times k$  applied to input  $I$ :

$$G[i, j] = \sum_{m=0}^{k-1} \sum_{n=0}^{k-1} F[m, n] \cdot I[i + m, j + n] + b$$

### 5.3.2 Output Size Calculation

For input size  $W \times H$ , filter size  $F$ , padding  $P$ , stride  $S$ :

$$W_{\text{out}} = \frac{W - F + 2P}{S} + 1$$

In our CNN:  $W = 28$ ,  $F = 3$ ,  $P = 0$ ,  $S = 1 \rightarrow W_{\text{out}} = 26$

### 5.3.3 Pooling Operation

Max pooling selects the maximum value in each window:

$$P[i, j] = \max\{G[Si : Si + F, Sj : Sj + F]\}$$

where  $S$  is stride and  $F$  is pool size.

## 5.4 CNN Compilation and Training

```
1 # =====
2 # CNN COMPIRATION
3 # =====
4 cnn.compile(
5     optimizer='adam',
6     loss='categorical_crossentropy',
7     metrics=['accuracy']
8 )
9
10 # =====
11 # CNN TRAINING
12 # =====
13 history_cnn = cnn.fit(
14     X_train_cnn,          # 4D tensor for CNN
15     y_train_cat,
16     epochs=10,
17     batch_size=32,
18     validation_data=(X_test_cnn, y_test_cat),
19     verbose=1
20 )
21
22 # CNN training characteristics:
23 # 1. Slower per epoch due to convolutional operations
24 # 2. Better feature extraction for spatial data
25 # 3. Parameter sharing reduces overfitting risk
26 # 4. Dropout provides additional regularization
```

Listing 14: CNN Compilation and Training

## 5.5 CNN Training Output Analysis

```
Epoch 1/10: accuracy: 0.7595 loss: 0.7384 val_accuracy: 0.9718
Epoch 2/10: accuracy: 0.9573 loss: 0.1400 val_accuracy: 0.9824
Epoch 3/10: accuracy: 0.9711 loss: 0.0908 val_accuracy: 0.9832
Epoch 4/10: accuracy: 0.9764 loss: 0.0711 val_accuracy: 0.9857
```

```

Epoch 5/10: accuracy: 0.9798 loss: 0.0636 val_accuracy: 0.9853
Epoch 6/10: accuracy: 0.9807 loss: 0.0554 val_accuracy: 0.9866
Epoch 7/10: accuracy: 0.9853 loss: 0.0455 val_accuracy: 0.9866
Epoch 8/10: accuracy: 0.9879 loss: 0.0376 val_accuracy: 0.9890
Epoch 9/10: accuracy: 0.9906 loss: 0.0312 val_accuracy: 0.9876
Epoch 10/10: accuracy: 0.9903 loss: 0.0309 val_accuracy: 0.9877

```

### 5.5.1 Key Observations

- **Highest accuracy:** 98.77% validation accuracy (best among all models)
- **Excellent generalization:** 98.77% validation vs 99.03% training
- **Fast convergence:** 97.18% accuracy in first epoch
- **Lowest loss:** 0.0309 training loss vs ANN's 0.0164
- **Stable learning:** Consistent improvement without oscillations

## 5.6 CNN Evaluation

```

1 # =====
2 # CNN EVALUATION
3 # =====
4 acc_cnn = cnn.evaluate(X_test_cnn, y_test_cat, verbose=0)[1]
5 print("\n" + "="*50)
6 print("CONVOLUTIONAL NEURAL NETWORK FINAL EVALUATION")
7 print("="*50)
8 print(f"Test Accuracy: {acc_cnn:.4f} ({acc_cnn*100:.2f}%)")

9
10 # Performance comparison with all models
11 print(f"\nCOMPREHENSIVE PERFORMANCE COMPARISON:")
12 print(f"Perceptron Accuracy: {acc_precp*100:.2f}%")
13 print(f"ANN Accuracy: {acc_ann*100:.2f}%")
14 print(f"CNN Accuracy: {acc_cnn*100:.2f}%")
15 print(f"\nImprovement over Perceptron: {((acc_cnn - acc_precp)*100:.2f}%
     ")
16 print(f"Improvement over ANN: {((acc_cnn - acc_ann)*100:.2f}%)"

```

Listing 15: CNN Performance Evaluation

## 6 Visualization and Analysis Functions

```

1 # =====
2 # VISUALIZATION FUNCTION
3 # =====
4 def plot_training(history, title):
5     """
6         Plots training and validation accuracy/loss.
7
8     Parameters:
9     -----
10    history : keras History object

```

```

11     Training history returned by model.fit()
12     title : str
13         Title for the plots
14     """
15     plt.figure(figsize=(12, 4))
16
17     # Plot Accuracy
18     plt.subplot(1, 2, 1)
19     plt.plot(history.history['accuracy'], label="Train", marker='o')
20     plt.plot(history.history['val_accuracy'], label="Validation",
21             marker='s')
22     plt.title(f"{title} - Accuracy")
23     plt.xlabel('Epoch')
24     plt.ylabel('Accuracy')
25     plt.legend()
26     plt.grid(True, alpha=0.3)
27
28     # Plot Loss
29     plt.subplot(1, 2, 2)
30     plt.plot(history.history['loss'], label="Train", marker='o')
31     plt.plot(history.history['val_loss'], label="Validation", marker='s')
32     plt.title(f"{title} - Loss")
33     plt.xlabel('Epoch')
34     plt.ylabel('Loss')
35     plt.legend()
36     plt.grid(True, alpha=0.3)
37
38     plt.tight_layout()
39     plt.show()
40
41 # Usage examples:
42 plot_training(history_precp, "Perceptron")
43 plot_training(history_ann, "Artificial Neural Network")
44 plot_training(history_cnn, "Convolutional Neural Network")
45
46 # =====
47 # CONFUSION MATRIX VISUALIZATION
48 # =====
49 def plot_confusion_matrix(model, x_test, y_test, model_name):
50     """
51         Plots confusion matrix for model predictions.
52     """
53     # Get predictions
54     predictions = model.predict(x_test)
55     predicted_classes = np.argmax(predictions, axis=1)
56     true_classes = np.argmax(y_test, axis=1)
57
58     # Calculate confusion matrix
59     cm = confusion_matrix(true_classes, predicted_classes)
60
61     # Plot
62     plt.figure(figsize=(10, 8))
63     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
64                 xticklabels=range(10), yticklabels=range(10))
65     plt.title(f'Confusion Matrix - {model_name}')
66     plt.ylabel('True Label')
67     plt.xlabel('Predicted Label')

```

```

67     plt.show()
68
69     # Calculate per-class accuracy
70     class_accuracy = cm.diagonal() / cm.sum(axis=1)
71     print(f"\nPer-class accuracy for {model_name}:")
72     for i, acc in enumerate(class_accuracy):
73         print(f"Digit {i}: {acc:.2%}")
74
75     # Usage:
76     # plot_confusion_matrix(Perceptron, x_test_img, y_test_cat, "Perceptron")
77     # plot_confusion_matrix(ann, x_test_img, y_test_cat, "ANN")
78     # plot_confusion_matrix(cnn, X_test_cnn, y_test_cat, "CNN")

```

Listing 16: Training Visualization Function

## 7 Comprehensive Analysis and Comparison

### 7.1 Parameter Count Analysis

Model	Parameters	Calculation	Efficiency
Perceptron	7,850	$(784 \times 10) + 10$	0.0115 acc/para
ANN	109,386	$(784 \times 128 + 128) + (128 \times 64 + 64) + (64 \times 10 + 10)$	0.000877 acc/par
CNN	225,034	$320 + 18,496 + 204,928 + 1,290$	0.000439 acc/par

Table 2: Parameter Efficiency Analysis

### 7.2 Performance Metrics Summary

Model	Train Acc	Val Acc	Train Loss	Val Loss
Perceptron	89.76%	90.32%	0.3759	0.3613
ANN	99.44%	95.98%	0.0164	0.1620
CNN	99.03%	98.77%	0.0309	0.0395

Table 3: Complete Performance Comparison

### 7.3 Training Time Analysis

Model	Time/Epoch	Total Time	Accuracy Gain/Time
Perceptron	2-3 seconds	20-30 seconds	3.01%/second
ANN	3-5 seconds	30-50 seconds	0.192%/second
CNN	18-22 seconds	180-220 seconds	0.0384%/second

Table 4: Training Efficiency Analysis

## 7.4 Error Analysis by Digit

- **Perceptron Errors:** Frequently confuses:
  - 3 vs 8 (similar rounded shapes)
  - 4 vs 9 (similar top structure)
  - 5 vs 6 (similar curves)
- **ANN Improvements:** Better at distinguishing:
  - 3 vs 8 (learns subtle curvature differences)
  - 4 vs 9 (recognizes open/closed loops)
  - Still struggles with poorly written digits
- **CNN Strengths:** Excellent at:
  - Recognizing digits regardless of position
  - Handling variations in writing style
  - Distinguishing subtle features
  - Generalizing to unseen variations

# 8 Key Insights and Learnings

## 8.1 Architectural Insights

1. **Model Complexity vs Performance:**
  - Perceptron: Simple but limited (90.32%)
  - ANN: Good balance (95.98%)
  - CNN: Complex but optimal (98.77%)
2. **Feature Learning Hierarchy:**
  - Perceptron: No feature learning, linear combination
  - ANN: Learns features but ignores spatial relationships
  - CNN: Hierarchical feature learning with spatial awareness
3. **Regularization Needs:**
  - Perceptron: Minimal overfitting due to simplicity
  - ANN: Shows overfitting (needs dropout/L2 regularization)
  - CNN: Dropout effectively prevents overfitting

## 8.2 Practical Recommendations

### 1. When to use each model:

- **Perceptron:** Simple binary classification, linearly separable data, limited resources
- **ANN:** Tabular data, moderate complexity, need for interpretability
- **CNN:** Image/video data, spatial patterns, maximum accuracy requirement

### 2. Training recommendations:

- Start with simple models, then increase complexity
- Use Adam optimizer for faster convergence
- Implement early stopping to prevent overfitting
- Use data augmentation for image data

### 3. Performance optimization:

- Batch size: 32-128 for good balance
- Learning rate: 0.001 for Adam, adjust based on progress
- Epochs: Monitor validation loss for early stopping
- Regularization: Use dropout (0.2-0.5) for complex models

## 8.3 Limitations and Future Improvements

### 1. Current Limitations:

- MNIST is relatively simple
- Models not tested on noisy/varied data
- No hyperparameter tuning performed
- Limited to grayscale images

### 2. Future Enhancements:

- Add batch normalization layers
- Implement data augmentation
- Try different architectures (ResNet, VGG)
- Add attention mechanisms
- Implement transfer learning

### 3. Advanced Techniques to Explore:

- Learning rate scheduling
- Gradient clipping
- Ensemble methods
- Bayesian optimization for hyperparameters
- Explainable AI techniques

## 9 Conclusion

### 9.1 Key Takeaways

1. **Architecture matters:** CNN significantly outperforms simpler models for image data
2. **Feature learning:** Hierarchical feature extraction in CNNs provides substantial benefits
3. **Computational trade-off:** More complex models require more resources but deliver better results
4. **Regularization is crucial:** Dropout effectively prevents overfitting in deep networks
5. **Optimizer choice:** Adam provides faster convergence than basic SGD

### 9.2 Final Performance Ranking

1. **1st:** CNN (98.77%) - Best for image classification
2. **2nd:** ANN (95.98%) - Good general-purpose model
3. **3rd:** Perceptron (90.32%) - Simple baseline model

### 9.3 Educational Value

This implementation provides:

- Hands-on experience with three fundamental architectures
- Understanding of neural network evolution
- Practical insights into model selection criteria
- Foundation for more advanced deep learning studies
- Template for comparative analysis methodology

### 9.4 Code Availability

The complete Jupyter notebook with all implementations, visualizations, and analysis is available for further study and experimentation. This serves as a comprehensive reference for understanding neural network architectures and their practical applications.